



Importing necessary libraries and packages

```
In [107... import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.decomposition import PCA
from collections import Counter
from sklearn.model_selection import GridSearchCV
```

TASK 1: Data Loading and Exploration

```
In [2]: df = pd.read_csv('data/MNIST.csv')
df.head()
```

```
Out[2]:
```

	1x1	1x2	1x3	1x4	1x5	1x6	1x7	1x8	1x9	1x10	...	28x20	28x21	28x22
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0

5 rows × 785 columns

Check for missing values

```
In [3]: df.isna()
```

```
Out[3]:
```

	1x1	1x2	1x3	1x4	1x5	1x6	1x7	1x8	1x9	1x10	...	28x20	2
0	False	False	False	False	False	False	False	False	False	False	...	False	
1	False	False	False	False	False	False	False	False	False	False	...	False	
2	False	False	False	False	False	False	False	False	False	False	...	False	
3	False	False	False	False	False	False	False	False	False	False	...	False	
4	False	False	False	False	False	False	False	False	False	False	...	False	
...	...	...	...	...	...	...	...	...	...	...	...	...	
69995	False	False	False	False	False	False	False	False	False	False	...	False	
69996	False	False	False	False	False	False	False	False	False	False	...	False	
69997	False	False	False	False	False	False	False	False	False	False	...	False	
69998	False	False	False	False	False	False	False	False	False	False	...	False	
69999	False	False	False	False	False	False	False	False	False	False	...	False	

70000 rows × 785 columns

```
In [4]: df.isna().sum().sum()
```

```
Out[4]: np.int64(0)
```

```
In [5]: df.isna().values.any()
```

```
Out[5]: np.False_
```

Class Distribution and number of samples

```
In [55]: df.shape[::]
```


```
Out[55]: (70000, 785)
```

```
In [32]: label_counts = df['label'].value_counts().sort_index()

plt.figure(figsize=(10,6))
plt.bar(label_counts.index, label_counts.values)

for i, v in enumerate(label_counts.values):
    plt.text(i, v+50, str(v), ha='center', fontsize=10)

plt.xlabel("Digit")
plt.ylabel("Frequency")
plt.title("Distribution of digit labels in MNIST data")
plt.xticks(range(10))
plt.show()
```



From the distribution of digit labels, it is visible that there is no significant class imbalance thereby providing accuracy to be a suitable evaluation metric.

Displaying sample images with labels


```
In [ ]: pixels = df.iloc[:, :-1] #separating pixel values and target labels
labels = df.iloc[:, -1]

sample_pixels = pixels.iloc[:10]
sample_labels = labels.iloc[:10]

plt.figure(figsize=(10,6))
for i in range(10):
    image = sample_pixels.iloc[i].values.reshape(28,28) #reshaping for image
    label = sample_labels.iloc[i]

    plt.subplot(2,5,i+1)
    plt.imshow(image, cmap='gray')
    plt.title(f"Label: {label}")
    plt.axis("off")

plt.tight_layout()
plt.show()
```



TASK 2: Data Preprocessing

Separating features and target column

```
In [72]: X = df.iloc[:, :-1]
y = df.iloc[:, -1]
print(X.shape, y.shape)
```

(70000, 784) (70000,)

Normalizing feature/pixel values

```
In [172]: X = X/255.0 #dividing by 255 as 0-255 is the range of pixel values in MNIST
```

Splitting dataset into training and testing sets (80:20 split)

```
In [77]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

(56000, 784) (14000, 784) (56000,) (14000,)

TASK 3: Model Implementation

Implementing KNN algorithm from scratch

```
In [83]: class KNN:
```

```

def __init__(self, k=3):
    """
    Number of K nearest neighbors. Default = 3
    """
    self.k = k

def fit(self, X, y):
    """
    Storing training data.
    """
    self.X_train = X.values if hasattr(X, "values") else X
    self.y_train = y.values if hasattr(y, "values") else y

def euclidian_distance(self, x1, x2):
    """
    Computing euclian distance between two points.
    """
    return np.sqrt(np.sum((x1-x2)**2))

def predict(self, X):
    """
    To predict labels for input samples X.
    """
    X = X.values if hasattr(X, "values") else X
    predictions = []

    for x in X:
        distances = [self.euclidian_distance(x, x_train) for x_train in self.X_train]
        k_indices = np.argsort(distances)[:self.k]
        k_labels = self.y_train[k_indices]
        common = Counter(k_labels).most_common(1)
        predictions.append(common[0][0])

    return np.array(predictions)

```

In [104... `from sklearn.utils import resample`


```
X_sub, y_sub = resample(X_train, y_train, n_samples=3000, random_state=42, str
```

```
In [98]: k_values = [1,3,5,7,9,11]

train_accuracies = []

for k in k_values:
    knn = KNN(k=k)
    knn.fit(X_sub, y_sub)
    preds = knn.predict(X_sub)
    acc = accuracy_score(y_sub, preds)
    train_accuracies.append(acc)
```

```
In [101]: plt.figure(figsize=(8, 5))
plt.plot(k_values, train_accuracies, marker='o')
plt.xlabel("Number of Neighbors (k)")
plt.ylabel("Training Accuracy")
plt.title("KNN Hyperparameter Tuning: Accuracy vs k")
plt.grid(True)
plt.xticks(range(12))
plt.show()
```



Although $k = 1$ yields the highest training accuracy, it is likely to overfit. A moderate value of $k = 5$ is chosen as a balance between model complexity and generalization.

```
In [141... best_knn = KNN(k=5)
best_knn.fit(X_train, y_train)
```

SVM without PCA

```
In [161... svm_without_pca = SVC(kernel='rbf')
param_grid = {
    'C' : [1, 5, 10],
    'gamma': ['scale', 0.01, 0.001]
}
```

```
In [162... X_svm, y_svm = resample(X_train, y_train, n_samples=3000, stratify=y_train)
```

```
In [163... grid_svm_without_pca = GridSearchCV(
    svm_without_pca,
    param_grid,
    cv=3,
    scoring='accuracy',
    n_jobs=-1
)

grid_svm_without_pca.fit(X_svm, y_svm)
```

```
Out[163... ► GridSearchCV ⓘ ?
► best_estimator_: SVC
    ► SVC ?
```

```
In [164... best_svm_without_pca = grid_svm_without_pca.best_estimator_
print('Best SVM parameters:', grid_svm_without_pca.best_params_)
```

Best SVM parameters: {'C': 10, 'gamma': 'scale'}

Feature/Dimensionality reduction using PCA for SVM

```
In [106... pca = PCA(n_components=0.95, random_state=42)

X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

print("Reduced dimensions:", X_train_pca.shape[1])
```

Reduced dimensions: 154

```
In [108... svm = SVC(kernel='rbf')
param_grid = {
    'C': [1, 5, 10],
    'gamma': ['scale', 0.01, 0.001]
}
```

Due to high computation time required for 56,000 samples for RBF-kernel training, a subset of 5,000 samples is used for training the final model on the full training set.

```
In [ ]: X_svm_sub, y_svm_sub = resample(X_train_pca,y_train,n_samples=5000,stratify=y_
```

```
In [142... grid_svm = GridSearchCV(  
    svm,  
    param_grid,  
    cv=3,  
    scoring='accuracy',  
    n_jobs=-1  
)  
  
grid_svm.fit(X_svm_sub, y_svm_sub)
```

```
Out[142...  
  ▶ GridSearchCV ⓘ ?  
  ▶ best_estimator_: SVC  
    ▶ SVC ?
```

```
In [ ]: best_svm = grid_svm.best_estimator_  
print('Best SVM parameters:', grid_svm.best_params_)
```

Best SVM parameters: {'C': 10, 'gamma': 'scale'}

Decision Tree Classifier

```
In [136... dt = DecisionTreeClassifier(random_state=42)
```

```
In [137... X_dt_sub, y_dt_sub = resample(X_train,y_train,n_samples=5000,stratify=y_train)
```

```
In [138... param_grid_dt = {  
    'max_depth': [10, 20, 30, None],  
    'min_samples_split': [2, 5, 10]  
}
```

```
In [139... grid_dt = GridSearchCV(  
    dt,  
    param_grid_dt,  
    cv=3,  
    scoring='accuracy',  
    n_jobs=-1  
)  
  
grid_dt.fit(X_dt_sub, y_dt_sub)
```


Out[139...

```
    ▸ GridSearchCV ⓘ ?  
    ▸ best_estimator_: DecisionTreeClassifier  
      ▸ DecisionTreeClassifier ?
```

In [140...

```
best_dt = grid_dt.best_estimator_  
print("Best Decision Tree parameters:", grid_dt.best_params_)
```

Best Decision Tree parameters: {'max_depth': 10, 'min_samples_split': 5}

Training all models on full training data

In [165...

```
best_knn = KNN(k=5)  
best_knn.fit(X_train, y_train)  
  
best_svm_without_pca = SVC(**grid_svm_without_pca.best_params_)  
best_svm_without_pca.fit(X_train, y_train)  
  
best_svm = SVC(**grid_svm.best_params_)  
best_svm.fit(X_train_pca, y_train)  
  
best_dt = DecisionTreeClassifier(**grid_dt.best_params_, random_state=42)  
best_dt.fit(X_train, y_train)
```

Out[165...

▼ DecisionTreeClassifier ⓘ ?		
Parameters		
 criterion		'gini'
 splitter		'best'
 max_depth		10
 min_samples_split		5
 min_samples_leaf		1
 min_weight_fraction_leaf		0.0
 max_features		None
 random_state		42
 max_leaf_nodes		None
 min_impurity_decrease		0.0
 class_weight		None
 ccp_alpha		0.0
 monotonic_cst		None

Task 4: Model Evaluation

```
In [173... X_test = X_test.reset_index(drop=True) #reset index to convert label indexin
y_test = y_test.reset_index(drop=True)
```

```
In [156... X_test_sub, y_test_sub = resample(
    X_test,
    y_test,
    n_samples=100,
    random_state=42,
    stratify=y_test
)
```

```
In [167... sub_idx = X_test_sub.index
X_test_pca_sub = X_test_pca[sub_idx]
```

```
In [ ]: y_pred_svm_sub = best_svm.predict(X_test_pca_sub)
```

KNN Evaluation

```
In [158... print("KNN Evaluation (Subset)")
```

```

y_pred_knn = best_knn.predict(X_test_sub)
acc_knn = accuracy_score(y_test_sub, y_pred_knn)
print(f"KNN Accuracy (subset): {acc_knn:.4f}")

cm_knn = confusion_matrix(y_test_sub, y_pred_knn)
print(cm_knn)

plt.figure(figsize=(6,5))
sns.heatmap(cm_knn, cmap="Blues", cbar=False)
plt.title("KNN Confusion Matrix (Subset)")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```


KNN Evaluation (Subset)

KNN Accuracy (subset): 0.9700

```

[[10  0  0  0  0  0  0  0  0  0]
 [ 0 11  0  0  0  0  0  0  0  0]
 [ 0  0 10  0  0  0  0  0  0  0]
 [ 1  0  0  9  0  0  0  0  0  0]
 [ 0  0  0  0  9  0  0  0  0  1]
 [ 0  0  0  0  0  9  0  0  0  0]
 [ 0  0  0  0  0  0 10  0  0  0]
 [ 0  0  0  0  0  0  0 10  0  0]
 [ 0  0  0  0  0  0  0  0 10  0]
 [ 0  0  0  0  1  0  0  0  0  9]]

```



SVM without PCA Evaluation

```
In [166... print("SVM Evaluation without PCA")

y_pred_svm = best_svm_without_pca.predict(X_test)
acc_svm = accuracy_score(y_test, y_pred_svm)
print(f"SVM Accuracy: {acc_svm:.4f}")

cm_svm_without_pca = confusion_matrix(y_test, y_pred_svm)
print(cm_svm_without_pca)


plt.figure(figsize=(6,5))
sns.heatmap(cm_svm_without_pca, cmap="Blues", cbar=False)
plt.title("SVM Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

SVM Evaluation without PCA

SVM Accuracy: 0.9834

```
[[1377 1 0 0 0 1 0 0 1 1]
 [ 0 1561 5 1 1 0 0 4 3 0]
 [ 5 1 1374 2 2 0 1 7 4 2]
 [ 1 2 6 1399 0 7 0 8 5 0]
 [ 2 3 1 0 1335 0 5 2 0 17]
 [ 2 0 0 6 4 1236 8 0 5 2]
 [ 9 2 1 0 2 2 1358 0 1 0]
 [ 2 3 7 1 6 0 0 1431 0 9]
 [ 2 3 4 4 1 7 2 1 1336 5]
 [ 1 6 0 5 7 3 0 6 2 1361]]
```

SVM Confusion Matrix



SVM with PCA Evaluation

```
In [153... print("SVM Evaluation (Full Test Set)")

y_pred_svm = best_svm.predict(X_test_pca)
acc_svm = accuracy_score(y_test, y_pred_svm)
print(f"SVM Accuracy: {acc_svm:.4f}")

cm_svm = confusion_matrix(y_test, y_pred_svm)
print(cm_svm)

plt.figure(figsize=(6,5))
```

```

sns.heatmap(cm_svm, cmap="Blues", cbar=False)
plt.title("SVM Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

SVM Evaluation (Full Test Set)


SVM Accuracy: 0.9859

```

[[1377  1  0  0  0  1  0  0  1  1]
 [  0 1560  8  0  1  0  0  4  2  0]
 [  3  0 1381  1  2  0  1  5  4  1]
 [  1  1  5 1402  0  6  0  6  6  1]
 [  1  4  1  0 1339  0  4  2  0 14]
 [  2  0  0  7  2 1238  8  0  3  3]
 [  5  2  0  0  2  1 1364  0  1  0]
 [  2  3  7  0  5  0  0 1434  0  8]
 [  2  3  1  4  0  4  3  1 1343  4]
 [  1  6  0  4  6  3  0  5  1 1365]]

```

SVM Confusion Matrix



In [154... print("Decision Tree Evaluation (Full Test Set)")

```

y_pred_dt = best_dt.predict(X_test)
acc_dt = accuracy_score(y_test, y_pred_dt)
print(f"Decision Tree Accuracy: {acc_dt:.4f}")

cm_dt = confusion_matrix(y_test, y_pred_dt)

```

```

print(cm_dt)

plt.figure(figsize=(6,5))
sns.heatmap(cm_dt, cmap="Blues", cbar=False)
plt.title("Decision Tree Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

Decision Tree Evaluation (Full Test Set)


Decision Tree Accuracy: 0.8611

```

[[1288   2   18   13   3   17   14   5   15   6]
 [  1 1508   26   9   9   4   2   8   7   1]
 [ 23   12 1185   24  24  17  28  18  48  19]
 [  7   9   52 1176   13  56  13  20  40  42]
 [  7   9   19   4 1167   12  11  12  20 104]
 [ 29   7   15  77  20 1016  35   6  27  31]
 [ 30  14   25  25  29  44 1179   8  17   4]
 [  4  14  35  31  12  17   2 1297  10  37]
 [ 18  28  40  40  25  44  20  15 1083  52]
 [ 10  13  21  28  73  30   6  22  32 1156]]

```

Decision Tree Confusion Matrix



Misclassifications with SVM (PCA)

```

In [170... y_pred_svm_sub = best_svm.predict(X_test_pca_sub)

```

```

mis_idx_svm = np.where(y_pred_svm_sub != y_test_sub)[0]

plt.figure(figsize=(10, 4))


for i, idx in enumerate(mis_idx_svm[:5]):
    img = X_test_sub.iloc[idx].values.reshape(28, 28)
    true_label = y_test_sub.iloc[idx]
    pred_label = y_pred_svm_sub[idx]

    plt.subplot(1, 5, i + 1)
    plt.imshow(img, cmap="gray")
    plt.title(f"T:{true_label}, P:{pred_label}")
    plt.axis("off")

plt.suptitle("Misclassified Digits - SVM")
plt.show()

```

Misclassified Digits — SVM



Misclassifications with Decision Tree

```

In [171]: y_pred_dt_sub = best_dt.predict(X_test_sub)
mis_idx_dt = np.where(y_pred_dt_sub != y_test_sub)[0]

plt.figure(figsize=(10, 4))

for i, idx in enumerate(mis_idx_dt[:5]):
    img = X_test_sub.iloc[idx].values.reshape(28, 28)
    true_label = y_test_sub.iloc[idx]
    pred_label = y_pred_dt_sub[idx]

    plt.subplot(1, 5, i + 1)
    plt.imshow(img, cmap="gray")
    plt.title(f"T:{true_label}, P:{pred_label}")
    plt.axis("off")


plt.suptitle("Misclassified Digits - Decision Tree")

```



```
plt.show()
```

Misclassified Digits — Decision Tree



Task 5: Reporting

OBSERVATIONS

- The accuracy of the KNN model is observed to be approximately **97%**. Since the model is implemented from scratch, a representative subset of the test data is used for evaluation to manage computational complexity. The optimal number of neighbors (**k = 5**) is selected to balance overfitting and generalization.
- The accuracy of SVM with PCA applied is slightly higher than basic SVM model without PCA (**98.34% vs 98.59%**). In both models, the optimal hyperparameters are observed to be '**C**' = **10** and '**gamma**' = '**scale**'.
- The decision tree model accuracy is concluded as **86.11%**. The optimal hyperparameters are **max_depth = 10** and **min_samples_split = 5**

MISCLASSIFICATION ANALYSIS

- As observed in the SVM model (with PCA), some of the misclassified digits are 4, 9, and 3. 4 is been mistakenly predicted as 9 and vice-versa. Similarly, in decision tree misclassifications, 7, 9, or 0 are mistakenly predicted as 5 or 3, while 3 is being predicted falsely as 8.
- The misclassifications primarily arise due to the loss of spatial information when images are flattened into one-dimensional feature vectors, making it difficult for classical machine learning models to capture subtle structural differences in handwritten digits.

SUGGESTIONS FOR IMPROVEMENT

- Techniques like image augmentation and using CNN models to learn spatial complexity can be used for improved inference.